

01_data_ingestion_03_data_transformation

April 8, 2024

```
[1]: import os
```

```
[2]: %pwd
```

```
[2]: 'D:\\Desktop\\Deep Learning\\Lab 7\\Deep-Learning-Model-Customization-and-  
Performance-Evaluation\\Research\\Cifar10'
```

```
[3]: os.chdir("../")  
os.chdir("../")
```

```
[4]: %pwd
```

```
[4]: 'D:\\Desktop\\Deep Learning\\Lab 7\\Deep-Learning-Model-Customization-and-  
Performance-Evaluation'
```

```
[5]: import logging  
from pathlib import Path  
logging.basicConfig(  
    # filename='extract_data.log',  
    level=logging.INFO,  
    format='%(asctime)s - %(levelname)s - %(message)s',  
    datefmt='%Y-%m-%d %H:%M:%S'  
)
```

```
[6]: print(Path(os.getcwd()))
```

```
D:\\Desktop\\Deep Learning\\Lab 7\\Deep-Learning-Model-Customization-and-  
Performance-Evaluation
```

```
[7]: import logging  
import os  
from dataclasses import dataclass  
from pathlib import Path  
import numpy as np  
import pandas as pd  
from sklearn.model_selection import train_test_split  
import tensorflow as tf  
from tensorflow.keras.datasets import cifar10  
from tensorflow.keras.utils import to_categorical
```

```

@dataclass
class DataTransformationConfig:
    root_dir: Path
    X_train_file: Path
    y_train_file: Path
    X_test_file: Path
    y_test_file: Path

class ConfigurationManager:
    def __init__(self):
        self.root_dir = Path(os.getcwd())
        self.X_train_file = self.root_dir / "Dataset/Modeltraining/Cifar10/
↪X_train.npy"
        self.y_train_file = self.root_dir / "Dataset/Modeltraining/Cifar10/
↪y_train.npy"
        self.X_test_file = self.root_dir / "Dataset/Modeltraining/Cifar10/
↪X_test.npy"
        self.y_test_file = self.root_dir / "Dataset/Modeltraining/Cifar10/
↪y_test.npy"

    def get_data_transformation_config(self) -> DataTransformationConfig:
        return DataTransformationConfig(
            root_dir=self.root_dir,
            X_train_file=self.X_train_file,
            y_train_file=self.y_train_file,
            X_test_file=self.X_test_file,
            y_test_file=self.y_test_file
        )

class DataTransformation:
    def __init__(self, config: DataTransformationConfig):
        self.config = config

    def ensure_directories_exist(self):
        """Ensure that directories for all file paths exist."""
        for path in [self.config.X_train_file, self.config.y_train_file, self.
↪config.X_test_file, self.config.y_test_file]:
            path.parent.mkdir(parents=True, exist_ok=True)

    def ensure_uniform_class_distribution(self, X, y):
        unique, counts = np.unique(y, return_counts=True)
        min_count = np.min(counts)
        avg_count = np.mean(counts)

```

```

# Logging the details about class distributions
logging.info(f"Number of unique classes in y: {len(unique)}")
logging.info(f"Unique classes in y: {unique}")
logging.info(f"Counts of each class in y: {counts}")
logging.info(f"Average frequency of all classes: {avg_count}")
logging.info(f"Minimum frequency among classes: {min_count}")

X_list = []
y_list = []
for class_value in unique:
    class_indices = np.where(y == class_value)[0]
    np.random.shuffle(class_indices)
    selected_indices = class_indices[:min_count]
    X_list.append(X[selected_indices, :])
    y_list.append(y[selected_indices])

    # Logging details about data removal for balancing
    logging.info(f"Class {class_value} reduced to {min_count} instances,
↳for balancing.")

X_balanced = np.concatenate(X_list, axis=0)
y_balanced = np.concatenate(y_list, axis=0)

# Shuffling the data to mix the classes
indices = np.arange(X_balanced.shape[0])
np.random.shuffle(indices)
X_balanced = X_balanced[indices]
y_balanced = y_balanced[indices]

# Logging final dataset size after balancing
logging.info(f"Number of rows in the balanced dataset: {X_balanced.
↳shape[0]}")

return X_balanced, y_balanced

def train_test_splitting(self):

    self.ensure_directories_exist()

    # Ensure the directory for the status file exists
    self.config.root_dir.mkdir(parents=True, exist_ok=True) # This line
↳ensures that the root directory exists

    # Load and preprocess CIFAR-10 dataset
    logging.info("Loading and preprocessing CIFAR-10 dataset...")
    (X_train, y_train), (X_test, y_test) = cifar10.load_data()
    X_train, X_test = X_train / 255.0, X_test / 255.0

```

```

y_train, y_test = to_categorical(y_train, 10), to_categorical(y_test, 10)

# Combine X_train and X_test to create X
X = np.concatenate((X_train, X_test), axis=0)

# Combine y_train and y_test to create y
y = np.concatenate((y_train, y_test), axis=0)

# Ensure uniform class distribution
X_balanced, y_balanced = self.ensure_uniform_class_distribution(X, y)

# Split the dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X_balanced, y_balanced, test_size=0.2, random_state=42)

# Save the split data as numpy files
np.save(self.config.X_train_file, X_train)
logging.info(f"X_train data saved to {self.config.X_train_file}")
np.save(self.config.y_train_file, y_train)
logging.info(f"y_train data saved to {self.config.y_train_file}")
np.save(self.config.X_test_file, X_test)
logging.info(f"X_test data saved to {self.config.X_test_file}")
np.save(self.config.y_test_file, y_test)
logging.info(f"y_test data saved to {self.config.y_test_file}")

def main():
    config_manager = ConfigurationManager()
    data_transformation_config = config_manager.get_data_transformation_config()
    data_transformation = DataTransformation(data_transformation_config)
    data_transformation.train_test_splitting()

if __name__ == "__main__":
    main()

```

2024-04-08 14:38:12 - WARNING - From D:\Desktop\Deep Learning\Lab 2\MNSIT-MLPClassifier\venv\lib\site-packages\keras\src\losses.py:2976: The name tf.losses.sparse_softmax_cross_entropy is deprecated. Please use tf.compat.v1.losses.sparse_softmax_cross_entropy instead.

2024-04-08 14:38:12 - INFO - Loading and preprocessing CIFAR-10 dataset...

2024-04-08 14:38:13 - INFO - Number of unique classes in y: 2

2024-04-08 14:38:13 - INFO - Unique classes in y: [0. 1.]

2024-04-08 14:38:13 - INFO - Counts of each class in y: [540000 60000]

2024-04-08 14:38:13 - INFO - Average frequency of all classes: 300000.0

2024-04-08 14:38:13 - INFO - Minimum frequency among classes: 60000

2024-04-08 14:38:13 - INFO - Class 0.0 reduced to 60000 instances for balancing.

```

2024-04-08 14:38:14 - INFO - Class 1.0 reduced to 60000 instances for balancing.
2024-04-08 14:38:15 - INFO - Number of rows in the balanced dataset: 120000
2024-04-08 14:38:31 - INFO - X_train data saved to D:\Desktop\Deep Learning\Lab
7\Deep-Learning-Model-Customization-and-Performance-
Evaluation\Dataset\Modeltraining\Cifar10\X_train.npy
2024-04-08 14:38:31 - INFO - y_train data saved to D:\Desktop\Deep Learning\Lab
7\Deep-Learning-Model-Customization-and-Performance-
Evaluation\Dataset\Modeltraining\Cifar10\y_train.npy
2024-04-08 14:38:35 - INFO - X_test data saved to D:\Desktop\Deep Learning\Lab
7\Deep-Learning-Model-Customization-and-Performance-
Evaluation\Dataset\Modeltraining\Cifar10\X_test.npy
2024-04-08 14:38:35 - INFO - y_test data saved to D:\Desktop\Deep Learning\Lab
7\Deep-Learning-Model-Customization-and-Performance-
Evaluation\Dataset\Modeltraining\Cifar10\y_test.npy

```

```

[13]: import pandas as pd
import numpy as np

# Specify the path to your numpy array file
numpy_file_path = 'Dataset/Modeltraining/Cifar10/X_train.npy'

# Load the numpy array file
data = np.load(numpy_file_path)

# Reshape the 4D array into a 2D array
num_images = data.shape[0]
num_pixels = data.shape[1] * data.shape[2] * data.shape[3]
data_reshaped = data.reshape(num_images, num_pixels)

# Convert the 2D array to a DataFrame
df = pd.DataFrame(data_reshaped)

# Display the DataFrame
df

```

```

[13]:
      0         1         2         3         4         5         6  \
0    0.560784  0.552941  0.533333  0.541176  0.529412  0.494118  0.478431
1    0.996078  0.996078  0.996078  0.992157  0.992157  0.992157  0.992157
2    0.423529  0.721569  0.952941  0.431373  0.717647  0.937255  0.443137
3    0.560784  0.580392  0.678431  0.572549  0.584314  0.682353  0.556863
4    0.247059  0.384314  0.701961  0.247059  0.384314  0.701961  0.250980
...      ...      ...      ...      ...      ...      ...      ...
95995  0.286275  0.321569  0.329412  0.254902  0.270588  0.266667  0.345098
95996  0.968627  0.862745  0.835294  0.960784  0.862745  0.835294  0.956863
95997  0.717647  0.611765  0.478431  0.713725  0.611765  0.478431  0.729412
95998  0.992157  1.000000  0.996078  0.996078  1.000000  0.996078  1.000000
95999  0.521569  0.596078  0.439216  0.494118  0.580392  0.419608  0.560784

```

	7	8	9	...	3062	3063	3064	\
0	0.466667	0.419608	0.498039	...	0.109804	0.486275	0.309804	
1	0.992157	0.992157	0.996078	...	0.086275	0.121569	0.149020	
2	0.721569	0.937255	0.443137	...	0.580392	0.525490	0.580392	
3	0.560784	0.662745	0.576471	...	0.470588	0.431373	0.384314	
4	0.388235	0.705882	0.258824	...	0.705882	0.235294	0.345098	
...	
95995	0.329412	0.329412	0.490196	...	0.231373	0.400000	0.356863	
95996	0.870588	0.839216	0.956863	...	0.752941	0.811765	0.729412	
95997	0.623529	0.494118	0.729412	...	0.345098	0.662745	0.545098	
95998	1.000000	1.000000	0.980392	...	1.000000	1.000000	1.000000	
95999	0.631373	0.478431	0.701961	...	0.129412	0.219608	0.243137	

	3065	3066	3067	3068	3069	3070	3071
0	0.117647	0.498039	0.317647	0.129412	0.498039	0.317647	0.125490
1	0.086275	0.078431	0.094118	0.043137	0.078431	0.047059	0.027451
2	0.572549	0.505882	0.580392	0.576471	0.501961	0.560784	0.560784
3	0.541176	0.415686	0.368627	0.525490	0.396078	0.349020	0.498039
4	0.701961	0.239216	0.352941	0.709804	0.239216	0.349020	0.709804
...
95995	0.203922	0.317647	0.298039	0.152941	0.266667	0.250980	0.109804
95996	0.749020	0.803922	0.725490	0.745098	0.807843	0.717647	0.729412
95997	0.345098	0.658824	0.537255	0.337255	0.654902	0.537255	0.337255
95998	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000
95999	0.156863	0.243137	0.262745	0.184314	0.313725	0.337255	0.243137

[96000 rows x 3072 columns]

```
[25]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt

# Specify the path to your numpy array file
np_file_path = 'Dataset/Modeltraining/Cifar10/X_train.npy'

# Load the numpy array file
data = np.load(np_file_path)

# Reshape the 4D array into a 2D array
num_images = data.shape[0]
num_pixels = data.shape[1] * data.shape[2] * data.shape[3]
data_reshaped = data.reshape(num_images, num_pixels)

# Convert the 2D array to a DataFrame
df = pd.DataFrame(data_reshaped)
```

```

# Select a random row number (change this as per your requirement)
row_number = 10

# Inverse transform the scaled row to get the original image
original_image = df.iloc[row_number].values.reshape(32, 32, 3)

# Display the image
plt.imshow(original_image)
plt.title(f"row_number: {row_number}") # Display the row number as the label
plt.axis('off')
plt.show()

```



```

[20]: import pandas as pd
import numpy as np

# Specify the path to your numpy array file
numpy_file_path = 'Dataset/Modeltraining/Cifar10/y_train.npy'

# Load the numpy array file
data = np.load(numpy_file_path)

# Reshape the 1D array into a 2D array with one column

```

```
data_resaped = data.reshape(-1, 1)

# Convert the 2D array to a DataFrame
df = pd.DataFrame(data_resaped, columns=['Label'])

# Display the DataFrame
df.head()
```

```
[20]:
```

	Label
0	0.0
1	0.0
2	0.0
3	1.0
4	0.0

```
[ ]:
```